

# **COMP2411**

# **REVIEW**

\* Final focus

File Organization and Indexing

Query Optimization

Transaction and Concurrency Control

# Lecture 01 Introduction to Database System

## Motivation

1. information about particular enterprise
2. file-processing system
3. disadvantages of traditional file-processing systems

**DB** non-redundant, persistent collection of logically related records/files that are structured to support various processing and retrieval needs.

**DBMS** set of software program for creating, storing, updating and accessing data of DB.

\* different b/t DBMS & other programming system

ability to manage persist data

DBMS: use conveniently, efficiently, robustly in retrieving & storing data

**Data Abstraction**: abstract view of data - Simplify interaction with system.

- hide detail how data is stored and manipulated

\* Tutorial level of abstraction:

physical/internal level : table

conceptual level : schema, what data are actually stored

View/external level : partial schema

1. The internal level has an **internal schema**, which describes the physical storage structure of the database. The internal schema uses a physical data model and describes the **complete details of data storage and access paths** for the database.

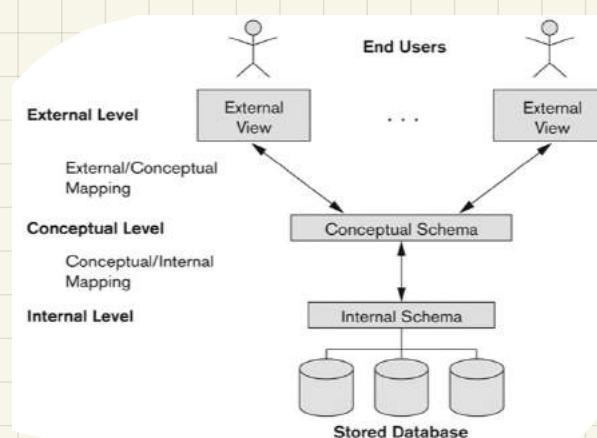
2. The **conceptual level** has a **conceptual schema**, which describes the structure of the whole database for a community of users. The conceptual schema hides the details of physical storage structures and **concentrates on describing entities, data types, relationships, user operations, and constraints**. Usually, a representational data model is used to describe the conceptual schema when a database system is implemented. This **implementation conceptual schema** is often based on a **conceptual schema design** in a high-level data model.

3. The **external or view level** includes a number of **external schemas** or user views. Each external schema describes the part of the database that a particular user group is interested in and hides the rest of the database from that user group. As in the previous level, each external schema is typically implemented using a representational data model, possibly based on an external schema design in a high-level data model.

**details of data storage & access path**

**describe entities/data type/relationships / user operation/constraints**

**describe parts of database that a particular user group is interested in and hides rest of the database from that user group**



• **Physical level**. The lowest level of abstraction describes *how* the data are actually stored. The physical level describes complex low-level data structures in detail.

• **Logical level**. The next-higher level of abstraction describes *what* data are stored in the database, and what relationships exist among those data. The logical level thus describes the entire database in terms of a small number of relatively simple structures. Although implementation of the simple structures at the logical level may involve complex physical-level structures, the user of the logical level does not need to be aware of this complexity. This is referred to as **physical data independence**. Database administrators, who must decide what information to keep in the database, use the logical level of abstraction.

• **View level**. The highest level of abstraction describes only part of the entire database. Even though the logical level uses simpler structures, complexity remains because of the variety of information stored in a large database. Many users of the database system do not need all this information; instead, they need to access only a part of the database. The view level of abstraction exists to simplify their interaction with the system. The system may provide many views for the same database.

**Data Model** conceptual tools for describing data

ER model (the entity-relationship model)

relational model

**Data Independence** ability to modify a schema definition in one level without affecting schema in next higher level.

physical data independence modify physical schema X alter conceptual schema X rewrite program

logical data independence modify conceptual schema X rewrite program

**Data Definition Language (DDL)**

compile to data dictionary (data about data).

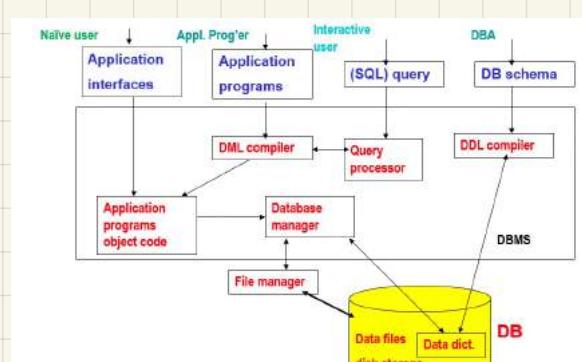
**Data Manipulation Language (DML)**

Query Language

**Database User** programmer, interactive users, naive users  
(DBA) DDL query language run program

**Database Administrator** central control over DB

**DB System Architecture**



## Course Objectives

- ER model: characterize relationships among entities
- Relational model: transform from ER diagram to tables
- SQL: language for writing queries
- Relational Algebra: logical way to represent queries
- Normal Forms: how to design good tables
- File Organization: provide file level structure to speed up query
- Query Optimization: transform queries into more efficient ones
- Transactions and Concurrency Control: handle concurrent operations and guarantee correctness of the database

## The Entity-Relationship Model (ER Model)

Entity distinguishable object with independent existence.

Entity Set set of entities of same type.

Attribute (property) information describing an entity.

Relationship an association among several entities.

Relationship Set set of relationships of the same type.

\* relationship can carry attributes

## Integrity Constraints

mapping cardinalities 1:1    1:M / N:1    M:N

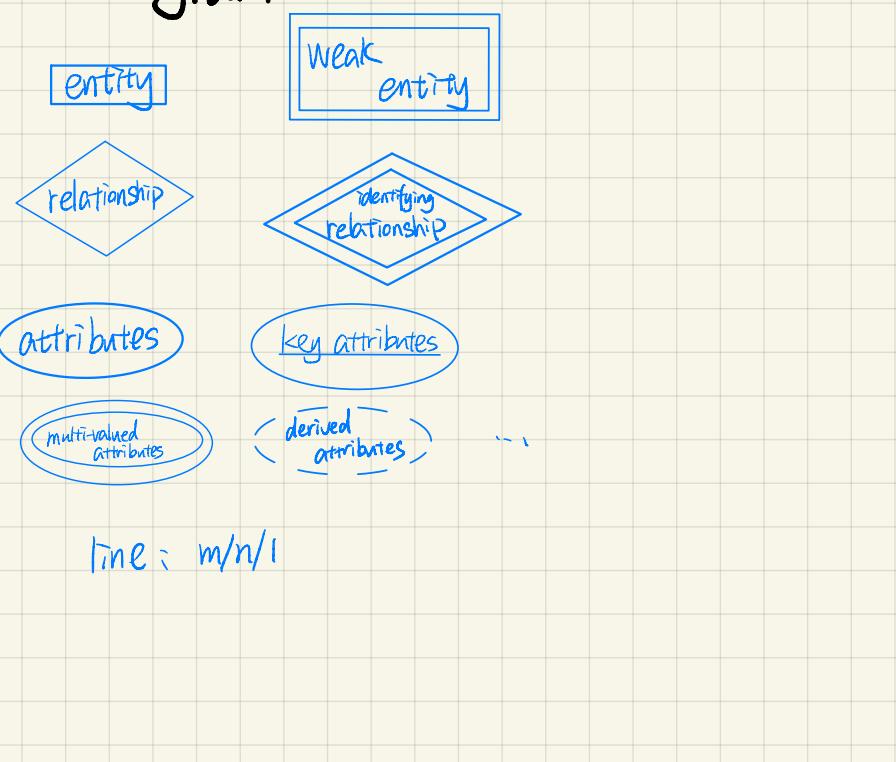
keys: distinguish individual entities or relationships

superkey set of attributes which identify uniquely an entity

candidate key minimal set of attributes which can identify unique an entity.

primary key candidate key chosen by DB designer to identify an entity.

## ER Diagram



## Summary of ER-Diagram Notation

Symbol	Meaning
rectangle	ENTITY TYPE
rectangle with diagonal line	WEAK ENTITY TYPE
diamond	RELATIONSHIP TYPE
double diamond	IDENTIFYING RELATIONSHIP TYPE
oval	ATTRIBUTE
oval with diagonal line	KEY ATTRIBUTE
oval with double line	MULTIVALUED ATTRIBUTE
oval with triple line	COMPOSITE ATTRIBUTE
oval with dashed line	DERIVED ATTRIBUTE
E <sub>1</sub> — R — E <sub>2</sub>	TOTAL PARTICIPATION OF E <sub>2</sub> IN R
E <sub>1</sub> — R <sup>N</sup> — E <sub>2</sub>	CARDINALITY RATIO 1:N FOR E <sub>1</sub> :E <sub>2</sub> IN R
R (min,max) — E	STRUCTURAL CONSTRAINT (min, max) ON PARTICIPATION OF E IN R

## Learning Objectives (ER model)

1. Explaining the meaning of Entity, Attribute and Relationship; able to find them on a ER diagram
2. On ER diagram, identify and explain the meaning of mapping cardinality, participation constraint
3. On ER diagram, able to explain every symbol covered in the lectures; able to transform ER-diagrams to relational tables
4. Able to use all the symbols to design a complete and correct ER diagram based on the user description of the application scenario

D: 1;    C: 1,2;    B: 1,2,3;    A:1,2,3,4

\*ER model

# Lecture 02 ER, Relational Data Model, SQL

## Attributes types

Simple single atomic value for attribute

Composite consist of several components, e.g. Name(Firstname, Lastname)

Multivalued multiple value for attribute, e.g. {color} → nested: example - { PreviousDegrees(College, Year, Degree, Field) }

## Relationship set/types

degree of relationship set number of participating entity sets/types binary, ternary, n-ary

\*binary: MANAGES & WORKS-ON

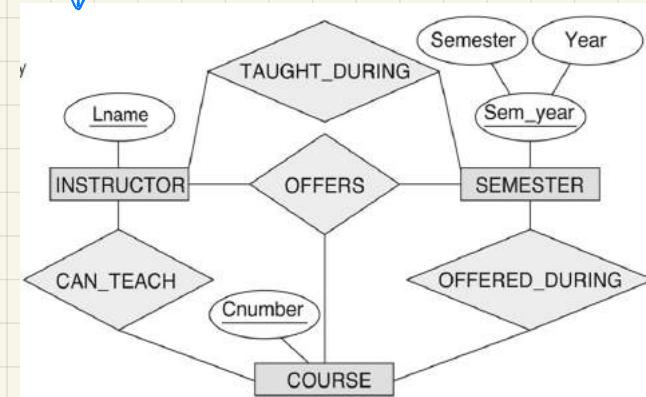
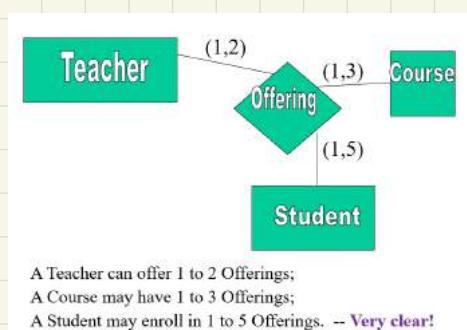
more than one relationship set exist with same entity sets

cardinality ratio: 1:1, 1:N, N:1, M:N

participation constraint partial participation - single line

total participation - double line

(min, max) notation



## Relational Data Model

### - Basic Structure

relations

data stored as tables(relations)

rows & columns (tuples & attributes)

attributes

domain

records

row/tuple in a relation

Schema

Customer			Parts		
cname	C#	address	P#	pname	cost

Orders			
order#	C#	P#	quantity
21	41256	301	15
26	41256	111	7
27	56437	301	11
32	23986	507	18
...			

DB instance  
rather  
dynamic

Customer			Parts		
cname	C#	address	P#	pname	cost
John	41256	8 Blue St., LA	301	widget	25,000
Mary	56437	6 Red Ave, SF	111	gadget	17,500
Joe	23986	12 Pink Rd, NY	507	screw	5,900

Orders			
order#	C#	P#	quantity
21	41256	301	15
26	41256	111	7
27	56437	301	11
32	23986	507	18
...			

## Characteristics of Relations

ordering of tuple in a relation r(R)

attributes in schema

value in a tuple

atomic / indivisible

table form X ordered

attributes ↗ order

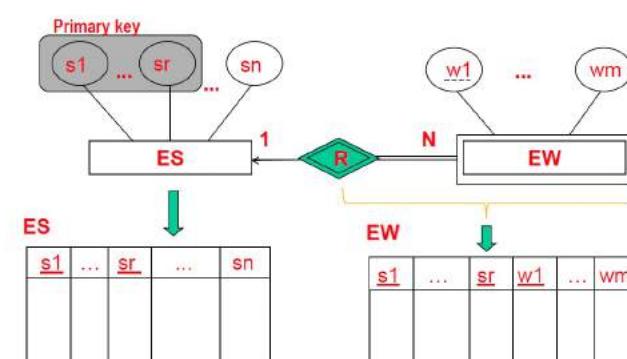
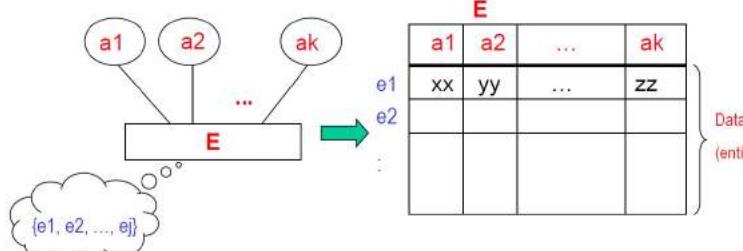
values ↗ order

\* null

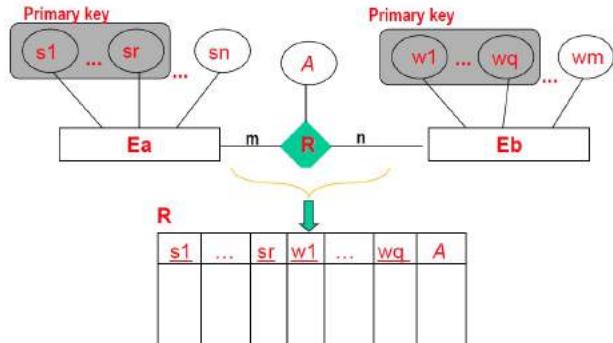
## ER Model vs. Relational Data Model

### Mapping ER Diagrams into Tables

#### ♦ Representation of (Strong) Entity Sets



#### ♦ Representation of M:N Relationship Sets



# Relational Query Language → relational algebra.

## SQL : Structed Query Language

### Learning Objectives (SQL)

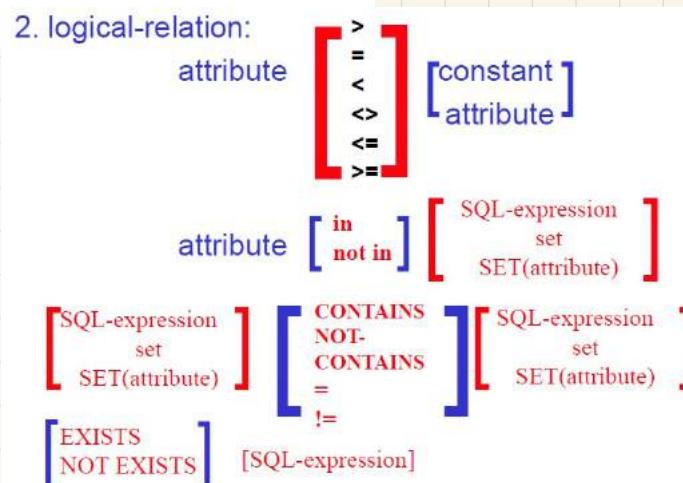
1. Able to explain the meaning of a basic SQL query (without aggregate functions and nested queries)
2. Given a requirement written in English, able to write a SQL query to achieve the requirement (without aggregate functions and nested queries)
3. Able to handle nested queries and views for 1 and 2
4. Able to handle aggregate functions for 1 and 2

D: 1; C: 1,2; B: 1,2,3; A:1,2,3,4

■ Basic syntax (structure) of SQL:

```
SELECT [DISTINCT] attribute1, ...
FROM table1 [name1], tabel2 [name2], ...
[WHERE requirement]
[GROUP BY (attribute) HAVING requirement];
```

1. requirement:

$$[\text{requirement} \underset{\text{logical-relation}}{\underset{\text{AND}}{\text{OR}}} \text{requirement}]$$


example

- Assuming the following relational schema:
  - Customer (cname, street, city)
  - Branch (bname, assets, b-city)
  - Borrow (bname, loan#, cname, amount)
  - Deposit (bname, acct#, cname, balance)
- ◆ Find all customers of the Sai Kong branch
 

```
(SELECT cname
       FROM Deposit
       WHERE bname = "Sai Kong")
UNION
(SELECT cname
       FROM Borrow
       WHERE bname = "Sai Kong");
```

- ◆ Find the name and city of all customers having a loan at the Sai Kong branch
 

```
SELECT DISTINCT Customer.cname, city
       FROM Borrow, Customer
       WHERE Borrow.cname = Customer.cname
          AND bname = "Sai Kong";
```
- ◆ Find the names of all customers whose street has the substring 'Main' included
 

```
SELECT cname
       FROM Customer
       WHERE street LIKE "%Main%";
```

(Note: if we use " \_\_\_ Main%", then it becomes a special case.)

◆ Find all customers who have an account at some branch in which Jones has an account

solution 1:

```
SELECT DISTINCT T cname
   FROM Deposit T
  WHERE T cname != "Jones"
    AND T bname IN (SELECT S bname
                    FROM Deposit S
                   WHERE S cname = "Jones");
```

solution 2:

```
SELECT DISTINCT T cname
   FROM Deposit S, Deposit T
  WHERE S cname = "Jones" AND S bname = T bname
    AND T cname != S cname;
```

*nested query*

◆ Find branches having greater assets than **all** branches in N.T.

solution 1:

```
SELECT X bname
   FROM Branch X
  WHERE NOT EXISTS (SELECT *
                      FROM Branch Y
                     WHERE Y b-city = "New Territory"
                       AND Y assets >= X assets);
```

*empty*

solution 2:

```
SELECT bname
   FROM Branch
  WHERE assets > ALL (SELECT assets
                      FROM Branch
                     WHERE b-city = "New Territory");
```

*all*

◆ Find names of all branches that have greater assets than **some** branch located in Kowloon

solution 1:

```
SELECT bname
   FROM Branch
  WHERE assets > SOME (SELECT assets
                      FROM Branch
                     WHERE b-city = "Kowloon");
```

solution 2:

```
SELECT X bname
   FROM Branch X, Branch Y
  WHERE X assets > Y assets
    AND Y b-city = "Kowloon";
```

◆ Find all customers who have a deposit account at **all** branches located in Kowloon

```
SELECT DISTINCT S cname
   FROM Deposit S
  WHERE (SELECT T bname
         FROM Deposit T
        WHERE S cname = T cname)
            CONTAINS
              (SELECT bname
                 FROM Branch
                WHERE b-city = "Kowloon");
```

*The set of branches he/she has a deposit account*

*(CONTAINS is dropped in most commercial languages)*

(Note: the last 3 examples show SQL supporting set comparisons)

The set of branches he/she has a deposit account

The set of branches in Kowloon

- Find all customers of Central branch who have an account there but no loan there

```
SELECT C.cname
FROM Customer C
WHERE NOT EXISTS (SELECT *
                   FROM Deposit D
                   WHERE D.cname = C.cname
                   AND D.bname = "Central")
AND NOT EXISTS (SELECT *
                 FROM Borrow B
                 WHERE B.cname = C.cname
                 AND B.bname = "Central");
```

- Find all customers who have a deposit account at ALL branches located in Kowloon

```
SELECT DISTINCT S.cname
FROM Deposit S
WHERE NOT EXISTS ((SELECT bname
                     FROM Branch
                     WHERE b-city = "Kowloon"))
MINUS
(SELECT T.bname
FROM Deposit T
WHERE S.cname = T.cname);
```

$S_1 - S_2 = \emptyset$

The set of branches located in KLN

The set of branches customer S holds some deposit accts

(Note: the last 2 examples show SQL can test whether a subquery has any tuples in its result, i.e., to test for empty tables.)

- List in alphabetic order all customers having a loan at branches in Kowloon

```
SELECT DISTINCT cname
FROM Borrow
WHERE b-city = "Kowloon"
ORDER BY cname;
```

By default, in ascending order.

ascending order

- List the entire borrow table in descending order of amount, and if several loans have the same amount, order them in ascending order by loan#

```
SELECT *
FROM Borrow
ORDER BY amount DESC, loan# ASC;
```

descending order

# SQL : Data Definition and Manipulation

## Table creation

```
CREATE TABLE table-name
(A1, D1,
A2, D2, ...
Ak, Dk);
```

## Table deletion

```
DROP TABLE table-name;
```

## Row/Tuple deletion

```
DELETE table-name;
```

## Table update

```
ALTER TABLE table-name
ADD Aj, Dj
```

(to add new attribute Aj with domain Dj to an existing table.)

## Views

- A view is a logical window of one or more tables
- in SQL, a view is defined using the CREATE VIEW command:

```
CREATE VIEW v-name AS <SQL-expression>;
```

## Deletion

### Syntax:

```
DELETE r
WHERE P
```

Meaning: those tuples t in relation r for which P(t) is true are deleted from r.

### Example1:

```
DELETE Deposit
WHERE b-name IN (SELECT b-name FROM Branch
                  WHERE b-city = "Laguna")
```

### Example2:

```
DELETE Deposit
WHERE balance < (SELECT AVG(balance) FROM Deposit)
      (Any possible anomalies/problems?)
```

## Update

### Syntax:

```
UPDATE r
SET clmn1 = expression
WHERE ...
```

Meaning: to change a value in one tuple (without changing all values in the tuple)

### Example1:

```
UPDATE Deposit
SET balance = balance * 1.05
      (to increase the payment by 5% to all accounts; it is applied to each tuple exactly once.)
```

### Example2:

```
UPDATE Deposit
SET balance = balance * 1.06 WHERE balance > 10000
UPDATE Deposit
SET balance = balance * 1.05 WHERE balance <= 10000
      (to increase the payment by 6% to all accounts with balance over $10000; all others receive 5% increase.)
```

E.g., suppose we define a view:

```
CREATE VIEW branch-city AS
SELECT bname, city
FROM Borrow, Customer
WHERE Borrow.cname = Customer.cname;
```

Now if we want to insert a tuple (Brighton, Shatin) through this view, it will cause:

(Brighton, null, null, null) => Borrow, and  
(null, null, Shatin) => Customer

Then what happens to the following query?

SELECT \* FROM branch-city;      null value comparison  
(would the tuple (Brighton, Shatin) be part of the answer?)      false

## Insertion

### Syntax:

```
INSERT INTO r VALUES (v1, v2, ..., vn);      or
```

INSERT INTO r

```
SELECT clmn_name1, clmn_name2, ..., clmn_namek
FROM s1, s2, ...
WHERE ...
```

### Example1:

```
INSERT INTO Deposit
SELECT b-name, loan#, c-name, 0
FROM Borrow WHERE b-name = "ClearWaterBay"
      (Meaning?)
```

# Null Value

Not	
T	F
U	U
F	T

AND	T	U	F
T	T	U	F
U	U	U	F
F	F	F	F

OR	T	U	F
T	T	T	T
U	T	U	U
F	T	U	F

- Find the average account balance at each branch:

```
SELECT b-name, AVG(balance)
FROM Deposit
GROUP BY b-name;
```

- Find those branches with the highest average balance:

```
SELECT b-name
FROM Deposit
GROUP BY b-name
HAVING AVG(balance) = SELECT MAX(AVG(balance))
FROM Deposit
GROUP BY b-name;
```

- If only interested in branches where average balance is > \$12000:

```
SELECT b-name, AVG(balance)
FROM Deposit
GROUP BY b-name
HAVING AVG(balance) > 12000;
```

- Find the average balance of all depositors who live in Laguna city and have at least 3 accounts:

```
SELECT AVG(balance)
FROM Deposit, Customer
WHERE Deposit.c-name = Customer.c-name AND city = "Laguna"
GROUP BY Deposit.c-name HAVING COUNT(DISTINCT acct#) >= 3;
```

By default, in ascending order.

ascending order

descending order

# Lecture 04 Relational Algebra

## Learning Objectives

1. Fully understand SELECT  $\sigma$  and PROJECT  $\pi$  operations.
2. Fully understand Set operations: UNION  $\cup$ , INTERSECTION  $\cap$ , DIFFERENCE  $-$ , CARTESIAN PRODUCT  $\times$ .
3. Fully understand JOIN operations.
4. Given a requirement written in English or SQL, you should be able to transform it into an expression in relational algebra; Given a relational algebra expression, you should be able to transform it into SQL and also explain its meaning

D: 1 C: 1,2 B: 1,2,3 A: 1,2,3,4

## SELECT $\sigma$

select  $\sigma_{c(R)}$  rows satisfy selection condition  $c$   
 Boolean expression

## Examples:

$\sigma_{DNO=4}(\text{EMPLOYEE})$

$\sigma_{\text{SALARY}>30000}(\text{EMPLOYEE})$

$\sigma_{(DNO=4 \text{ AND } \text{SALARY}>25000) \text{ OR } DNO=5}(\text{EMPLOYEE})$

## PROJECT $\pi$

keep certain attributes specified in an attribute list  $L$   
 $\pi_L(R)$  (columns)

Example:  $\pi_{FNAME,LNAME,SALARY}(\text{EMPLOYEE})$

Example:  $\pi_{SEX,SALARY}(\text{EMPLOYEE})$

Combine

Example: Retrieve the names and salaries of employees who work in department 4:

$\pi_{FNAME,LNAME,SALARY}(\sigma_{DNO=4}(\text{EMPLOYEE}))$

- Alternatively, we can specify explicit intermediate relations for each step:

$DEPT4\_EMPS \leftarrow \sigma_{DNO=4}(\text{EMPLOYEE})$

$R \leftarrow \pi_{FNAME,LNAME,SALARY}(DEPT4\_EMPS)$

## SET Operation

UNION

$R1 \cup R2$

INTERSECTION

$R1 \cap R2$

DIFFERENCE

$R1 - R2$

CARTESIAN PRODUCT  $R1 \times R2$

STUDENT	
Fn	Ln
Susan	Yao
Ramesh	Shah
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert

INSTRUCTOR	
Fname	Lname
John	Smith
Ricardo	Browne
Susan	Yao
Francis	Johnson
Ramesh	Shah

union	
Fn	Ln
Susan	Yao
Ramesh	Shah
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert
John	Smith
Ricardo	Browne
Francis	Johnson

intersect	
Fn	Ln
Susan	Yao
Ramesh	Shah

stu-instructor	
Fn	Ln
Johnny	Kohler
Barbara	Jones
Amy	Ford
Jimmy	Wang
Ernest	Gilbert

instructor-stu	
Fname	Lname
John	Smith
Ricardo	Browne
Francis	Johnson

As another example: Combine each female EMPLOYEE tuple with the DEPENDENTS tuple to get her dependent(s).

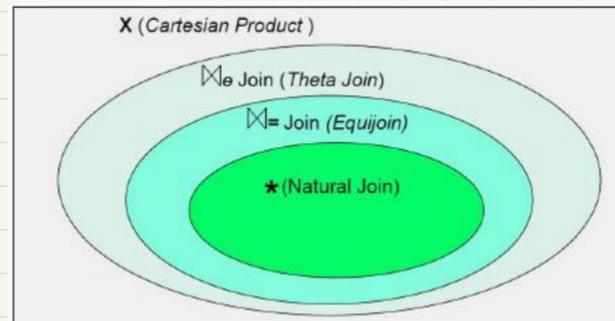
$\text{FemEMPNames} \leftarrow \pi_{FNAME,LNAME,SSN}(\sigma_{Sex=F}(\text{EMPLOYEE}))$

$\text{EMPDEPENDENTS} \leftarrow \text{FemEMPNames} \times \text{DEPENDENTS}$

$\text{ActualDEPENDENTS} \leftarrow \sigma_{SSN=ESSN}(\text{EMPDEPENDENTS})$

$\text{Result} \leftarrow \pi_{FNAME,LNAME,DEPENDENT\_NAME}(\text{ActualDEPENDENTS})$

## JOIN Operation



### Theta join

Cartesian product followed by a SELECT condition  $c$  ( $\theta \rightarrow$  join condition).

$R(A_1, A_2, \dots, A_m, B_1, B_2, \dots, B_n) \leftarrow R_1(A_1, A_2, \dots, A_m) \bowtie_c R_2(B_1, B_2, \dots, B_n)$

### Equijoin

join condition  $c$  (one or more equality comparison)

$C = (A_i=B_j) \text{ AND } \dots \text{ AND } (A_h=B_k); 1 \leq i, h \leq m, 1 \leq j, k \leq n$

$A_i - A_h \Rightarrow$  join attribute of  $R_1$   
 $B_j - B_k \Rightarrow$  join attribute of  $R_2$

### Natural Join

redundant join attributes of  $R_2$  is eliminated.

$R \leftarrow R_1 * (\text{join attributes of } R_1, \text{join attributes of } R_2)$

Example: Retrieve each EMPLOYEE's name and the name of the DEPARTMENT he/she works for:

$T \leftarrow \text{EMPLOYEE} *_{(DNO),(DNUMBER)} \text{DEPARTMENT}$

$\text{RESULT} \leftarrow \pi_{FNAME,LNAME,DNAME}(T)$

# LEC05 Integrity Constraints (ICs) & Normal Forms

## Learning Objectives

- Understand all Integrity Constraints covered in the lecture
  - Understand Functional Dependency, able to judge which set of attributes is a candidate key by looking at the given FDs
  - Understand Normalization, able to judge whether a table with a given set of FDs is 1NF, 2NF, 3NF or BCNF
  - Given an initial table together with a set of FDs, able to turn it step by step into a set of tables which is in BCNF
- D: 1 C: 1,2 B: 1,2,3 A: 1,2,3,4

## From Lecture 1

keys: distinguish individual entities or relationships  
superkey set of attributes which identify uniquely an entity  
candidate key minimal set of attributes which can identify unique an entity.  
primary key Candidate key chosen by DB designer to identify an entity.

## Key Constraints

Super key no two tuple have same value for superkey.

Candidate key minimal superkey (delete any one)

Primary key choice by DB designer to identify an entity  
underlined

## Domain/Entity Integrity Constraints

specify set of possible value (may be associate with an attribute)

Primary key  $\neq \text{null}$  \* other: may not allow null

## Referential Integrity Constraints (foreign key) involving 2 relations

specify a relationship among tuples in 2 relations

ensure a value appear in one relation also appear in another

\* subset dependency

## Functional Dependency (FD)

relation

Let R be a relation scheme, and  $\alpha \subseteq R$ ,  $\beta \subseteq R$  (ie,  $\alpha$  and  $\beta$  are sets of R's attributes). We say:

$$\alpha \rightarrow \beta$$

if in any legal relation instance r(R), for all pairs of tuples t1 and t2 in r, we have:

$$(t1[\alpha] = t2[\alpha]) \Rightarrow (t1[\beta] = t2[\beta])$$

## Inference Rules for FD

IR1 Reflexive  $Y \subseteq X \Rightarrow X \rightarrow Y$

IR2 Augmentation  $X \rightarrow Y \Rightarrow XZ \rightarrow YZ$  ( $XZ = X \cup Z$ )

IR3 Transitive  $\begin{matrix} X \rightarrow Y \\ Y \rightarrow Z \end{matrix} \Rightarrow X \rightarrow Z$

## Relative DB Design normalization theory \* lossless join

1NF if all underlying domain contains atomic value only  
problem: insert/delete/update anomalies.

$\downarrow$   
X represent certain info  
 $\downarrow$   
destroy all info  
 $\downarrow$   
inconsistency

2NF in 1NF and every non-key attribute is fully depend on any candidate key.

problem: insert / delete/update anomalies  
 $\downarrow$   
X insert if exists some info  
destroy all info  
 $\downarrow$   
potential inconsistency

3NF in 2NF and every non-key value is non-transitively depends on any candidate key.

problem: redundancy

BCNF multiple candidate keys (composite one), overlaps common attributes

E.g.,

C-Name: string of char (30)

Balance: Number (6,2)

```
CREATE TABLE Customer
(C-name Char(20) NOT NULL,
Street Char(30),
City Char(25),
PRIMARY KEY (C-name));
```

```
CREATE TABLE Branch
(B-name Char(20) NOT NULL,
Asset Integer,
B-City Char(20),
PRIMARY KEY (B-name));
```

```
CREATE TABLE Deposit
(B-name Char(20),
Acct# Char(10) NOT NULL,
C-name Char(20) NOT NULL,
Balance Integer,
PRIMARY KEY (Acct#, C-name),
FOREIGN KEY (B-name) REFERENCES Branch,
FOREIGN KEY (C-name) REFERENCES Customer);
```

## Equivalence of FDs

Two sets of FDs F and G are equivalent if:

- Every FD in F can be inferred from G, and
- Every FD in G can be inferred from F
- Hence, F and G are equivalent if  $F^+ = G^+$

Decomposition

$$X \rightarrow YZ \Rightarrow \begin{cases} X \rightarrow Y \\ X \rightarrow Z \end{cases}$$

Union

$$\begin{matrix} X \rightarrow Y \\ X \rightarrow Z \end{matrix} \Rightarrow X \rightarrow YZ$$

Pseudotransitivity

$$\begin{matrix} X \rightarrow Y \\ WY \rightarrow Z \end{matrix} \Rightarrow WX \rightarrow Z$$

e.g. FIRST(S#, Status, City, P#, Qty)

e.g. SECOND(S#, Status, City)  
SP(S#, P#, Qty)

NFs Defined Informally:

- 1<sup>st</sup> normal form
  - All attributes are of single atomic values which are dependent on the key
- 2<sup>nd</sup> normal form
  - All non-key attributes depend on the entire key
- 3<sup>rd</sup> normal form
  - All non-key attributes only depend on the key

# Lecture 07 File Organization & Indexing

## Learning Objectives

- File organization
    - ◆ Unordered file 1
    - ◆ Ordered file 2
    - ◆ Hash file 3
    - ◆ Dynamic and extendible hash 4
  
  - Indexes
    - ◆ Primary Indexes 1
    - ◆ Clustering Indexes 2
    - ◆ Secondary Indexes 3
    - ◆ Multilevel Indexes 4
    - ◆ Definitions of and operations on B-Trees and B<sup>+</sup>-Trees 3
- D: 1    C: 1,2    B: 1,2,3    A: 1,2,3,4    61

## File Operation

### \* File

#### Unordered File

insert new record → end, efficient

Search: linear search - read half block require sorting

#### Ordered File

sequential file (sorted by value)

Insertion: expensive - to correct position

Search: binary search log<sub>n</sub> efficient

#### Hashed File

external hashing

hash key ← one of file field

buckets ← file blocks

$i = h(K)$   $h$  = hashing function

Solving collision problem

- linear Probing

- Quadratic Probing

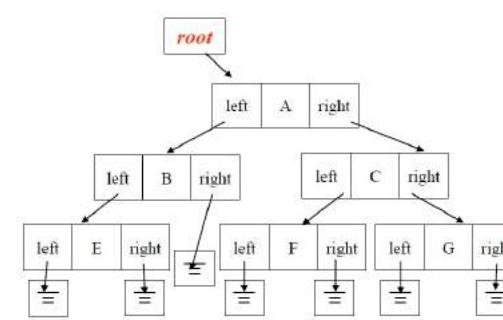
disadvantage: fixed bucket when records grows/shrinks

Search efficient.

Collision occurs when bucket is full

\* overflow

\* binary search with binary tree



```
class TreeNode {
    private:
        int info;
        TreeNode* left;
        TreeNode* right;
};

class Mytree {
    private:
        TreeNode* root;
}
```

\* assume  $h(x) = x \% D$

Homework

- 4) If a hash structure is used on a search key for which range queries are likely, what property should the hash functions have? [10 marks]

Answer:

The hash function should preserve the order of the data. That is, if  $h$  is the hash function, and  $x < y$ , then it implies that  $h(x) < h(y)$ .

### Dynamic and extendible hash

binary representation of hash value

dynamic hashing → binary tree

extendible hashing → array in  $2^d$  ( $d \rightarrow$  global depth)

### \* Type of Single-Level Indexes

#### Primary Indexes

on ordered data (key field)

fixed length size with two field primary key specific data block

dense index → search faster

→ more space for store index record

Sparse index → search slower

→ need less space

#### Clustering Indexes

on ordered data (non-key field)

record stored in index (not pointer)

not unique for each record

#### Secondary Indexes

ordered data (candidate key & non-key field)

non-ordering index

reduce mapping size

Homework when to use dense index?

- 2) When is it preferable to use a dense index rather than a sparse index? [8 marks]

Answer:

It is preferable to use a dense index rather than a sparse index when

- the file is not ordered on the indexed field (such as when the index is a secondary index) (4 marks)
- or, when the index file is small enough to fit in the main memory (4 marks)

Homework

- 3) What is the difference between primary index and a secondary index? [7 marks]

Answer:

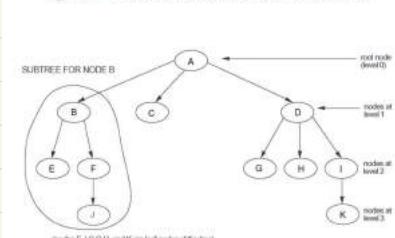
The difference between primary index and a secondary index lies in the facts that:

- The primary index is on the field which specifies the sequential order of the data file (4 marks)
- There can only be one primary index while many secondary indices (3 marks)

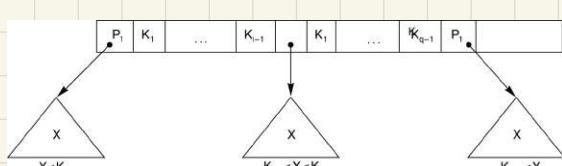
### Multilevel Indexes

index to index → second-level index

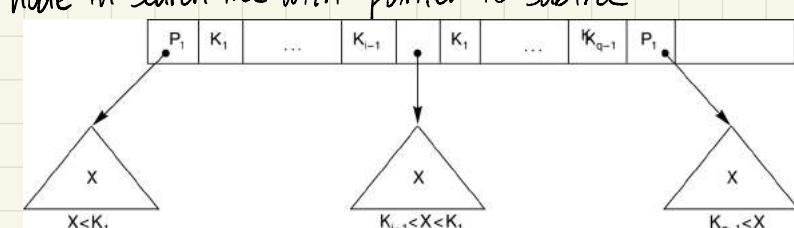
Figure 6.7 A tree data structure that shows an unbalanced tree.



search tree of order  $p=3$  (3 pointer)



node in search tree with pointer to subtree



# B Tree and B<sup>+</sup> Tree

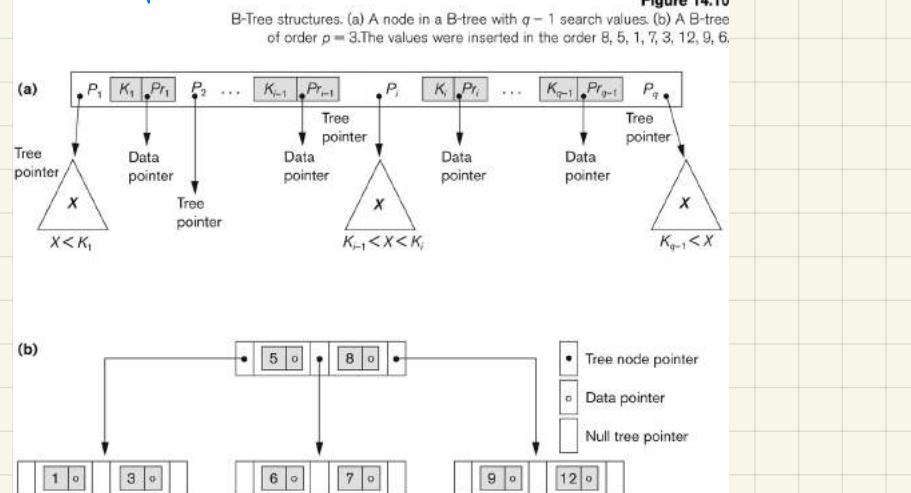
- leave space for each tree node
- each node ↔ a disk block
- node: half full & completely full

insertion: if node not full, add  
if node full, split to two nodes → other tree level

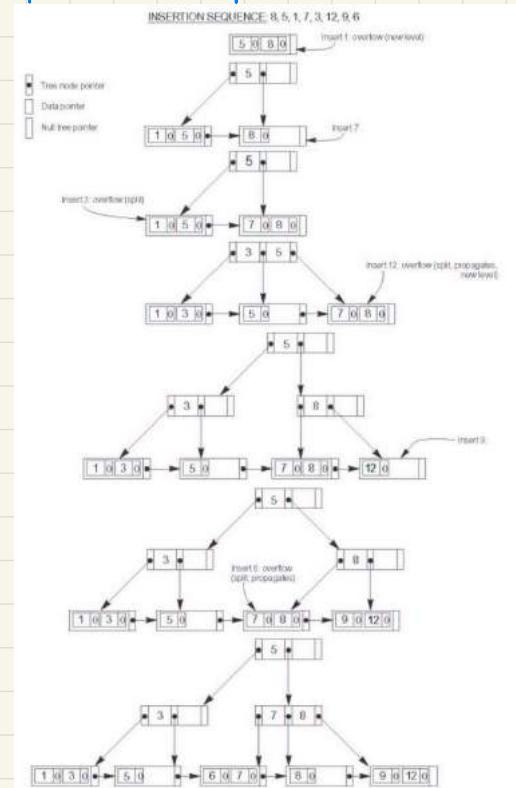
deletion if node X become less than half full, delete  
else → merge neighboring nodes

difference: B tree pointer to data records exist at all level of tree  
B<sup>+</sup> tree pointer to data records exist at leaf level

## Example B trees



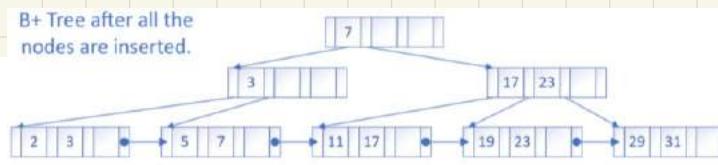
Example insertion of B<sup>+</sup> tree  $P=3$  Leaf = 2



## \* Tutorial

2. Construct a B<sup>+</sup> tree for the following set of key values:  
(2, 3, 5, 7, 11, 17, 19, 23, 29, 31)

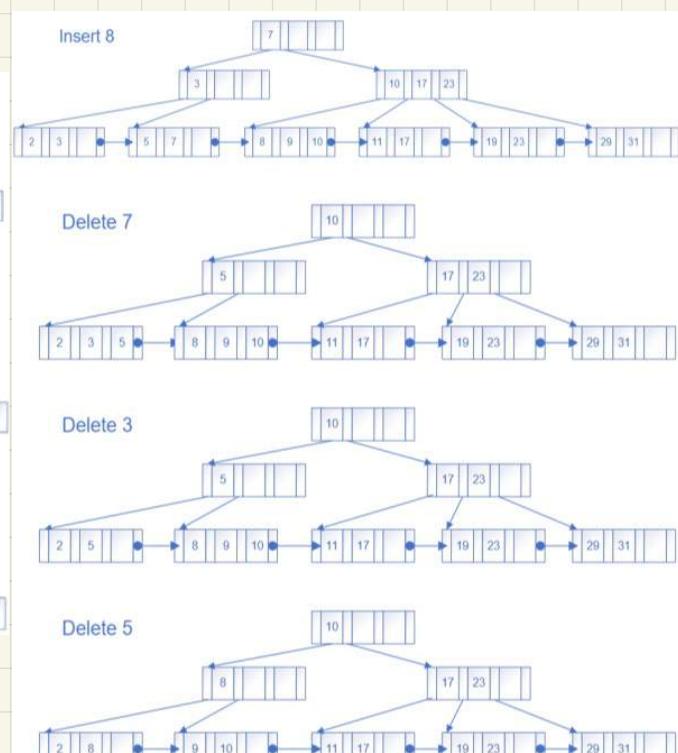
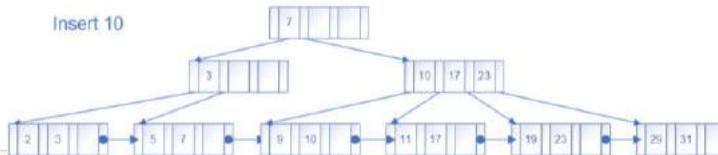
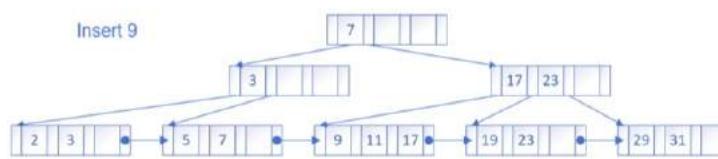
B<sup>+</sup> Tree after all the nodes are inserted.



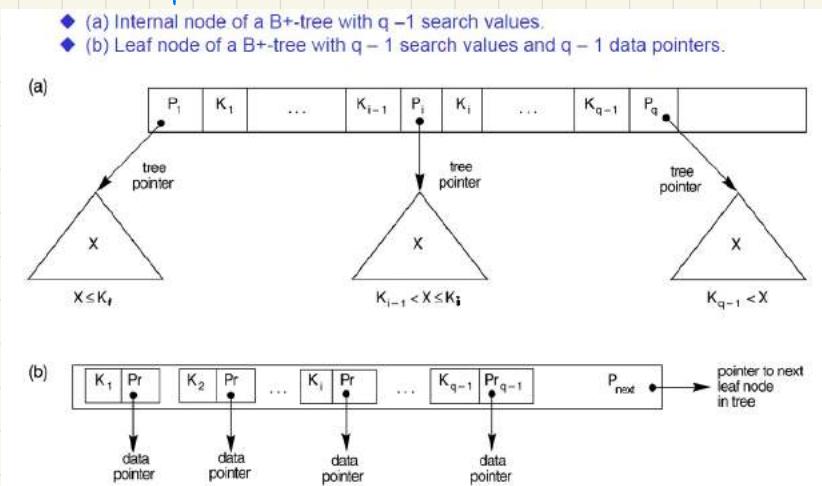
Assuming that the tree is initially empty, values are added in ascending order, and the order of the tree is 4 and the leaf order is 3.

For the tree derived from the above question, show the form of the tree after each of the following series of operations:

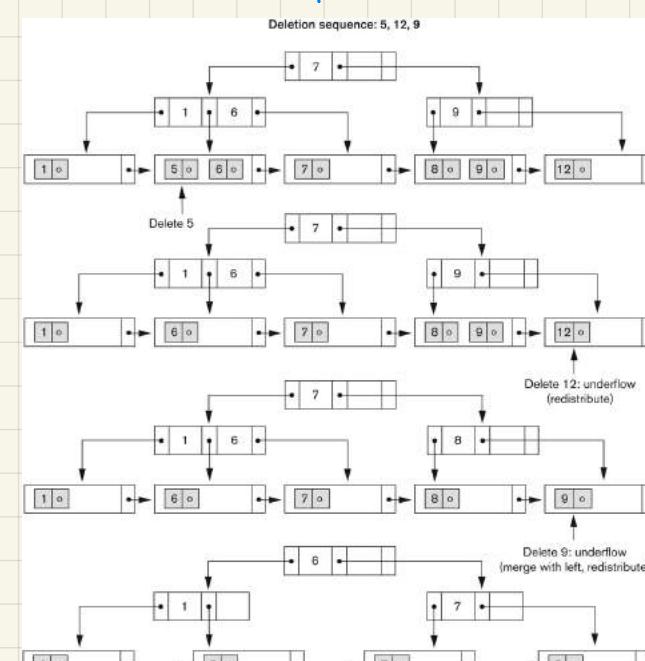
- Insert 9
- Insert 10
- Insert 8
- Delete 7
- Delete 3
- Delete 5



## Example B<sup>+</sup> trees



## Example Deletion of B<sup>+</sup> tree



# Lecture 09 Query Optimization

## Learning Objectives

- Different ways of implementing "Select". Able to select suitable ways when certain conditions are given.
- Different ways of implementing "Join". Able to select suitable ways when certain conditions are given; Able to compare different methods (which one is better) given certain condition.
- Be familiar with the ways to implement aggregate functions utilizing index structure.
- Understand how to improve the execution of queries, and be able to use those 12 heuristic rules

D: 1 C:1,2 B: 1,2,3 A: 1,2,3,4

## Query Optimization

- SQL  $\Rightarrow$  Relational Algebra

- Search Method for SELECT

- Linear search every record
- Binary search equality comparison (ordered file)
- Use a primary index/hash key to retrieve single record  
equality comparison on key attribute with a primary index/hash key
- use a primary index to retrieve multiple records  
 $> \geq < \leq$  on key field with primary index
- use clustering index to retrieve multiple records  
equality comparison on non-key attribute with clustering index
- use secondary index / B<sup>+</sup> tree  $> \geq < \leq$  (range queries)
- conjunctive selection
- conjunctive selection using complex index
- conjunctive selection by intersection

## Join Operation

- Nested (inner-outer) loop approach
- access structure to retrieve matching records
- Sort-merge
- hash join

## Aggregated Operation

{ table scan  
index  
max, min, sum, count, avg

## Heuristic Optimization

### Query Tree

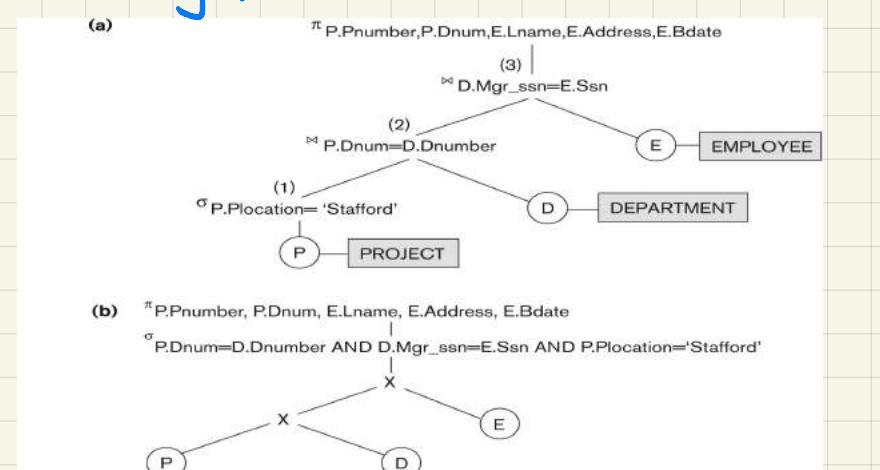


Figure 15.4  
Two query trees for the query Q2. (a) Query tree corresponding to the relational algebra expression for Q2. (b) Initial (canonical) query tree for SQL query Q2. (c) Query graph for Q2.

### Rules

- Cascade of  $\sigma$
- Commutative of  $\sigma$
- Cascade of  $\pi$
- Commuting with  $\pi$
- Commutativity of \* (or  $\bowtie$ )
- Commuting  $\sigma$  with \*
- Commutativity of set operation
- Commuting  $\pi$  with  $\bowtie$
- Associativity \*,  $\times$ ,  $\cup$ ,  $\cap$
- Commuting  $\sigma$  with set operation
- $\pi$  operation commuting with  $\cup$

$$\sigma_{c_1} \text{ AND } c_2 \text{ AND } \dots \text{ AND } c_n(R) = \sigma_{c_1}(\sigma_{c_2}(\dots(\sigma_{c_n}(R))\dots))$$

$$\sigma_{c_1}(\sigma_{c_2}(R)) = \sigma_{c_2}(\sigma_{c_1}(R))$$

$$\pi \text{ List1 } (\pi \text{ List2 } (\dots(\pi \text{ Listn}(R))\dots)) = \pi \text{ List1}(R)$$

$$\pi A_1, A_2, \dots, A_n(\sigma_c(R)) = \sigma_c(\pi A_1, A_2, \dots, A_n(R))$$

$$R * S = S * R$$

$$\sigma_c(R * S) = (\sigma_c(R)) * S$$

$$\cap$$

$$\cup$$

$$\pi_L(R \bowtie_c S) = (\pi_{A_1, \dots, A_n}(R)) \bowtie_c (\pi_{B_1, \dots, B_m}(S))$$

$$(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$$

$$\sigma_c(R \bowtie S) = (\sigma_c(R)) \bowtie (\sigma_c(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

$$\pi_L(R \bowtie S) = (\pi_L(R)) \bowtie (\pi_L(S))$$

$$\pi_L(R \cap S) = (\pi_L(R)) \cap (\pi_L(S))$$

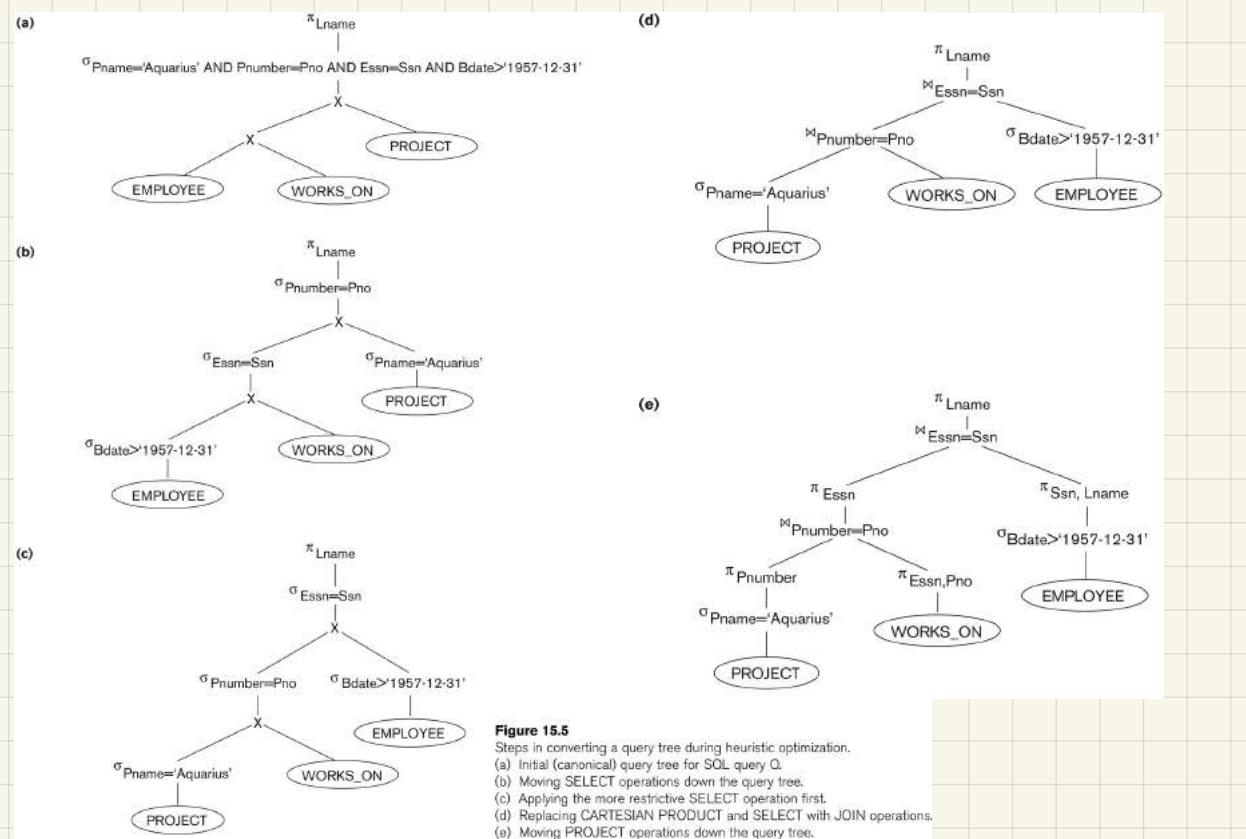
$$\pi_L(R \cup S) = (\pi_L(R)) \cup (\pi_L(S))$$

$$\pi_L(R - S) = (\pi_L(R)) - (\pi_L(S))$$

$$\pi_L(R \times S) = (\pi_L(R)) \times (\pi_L(S))$$

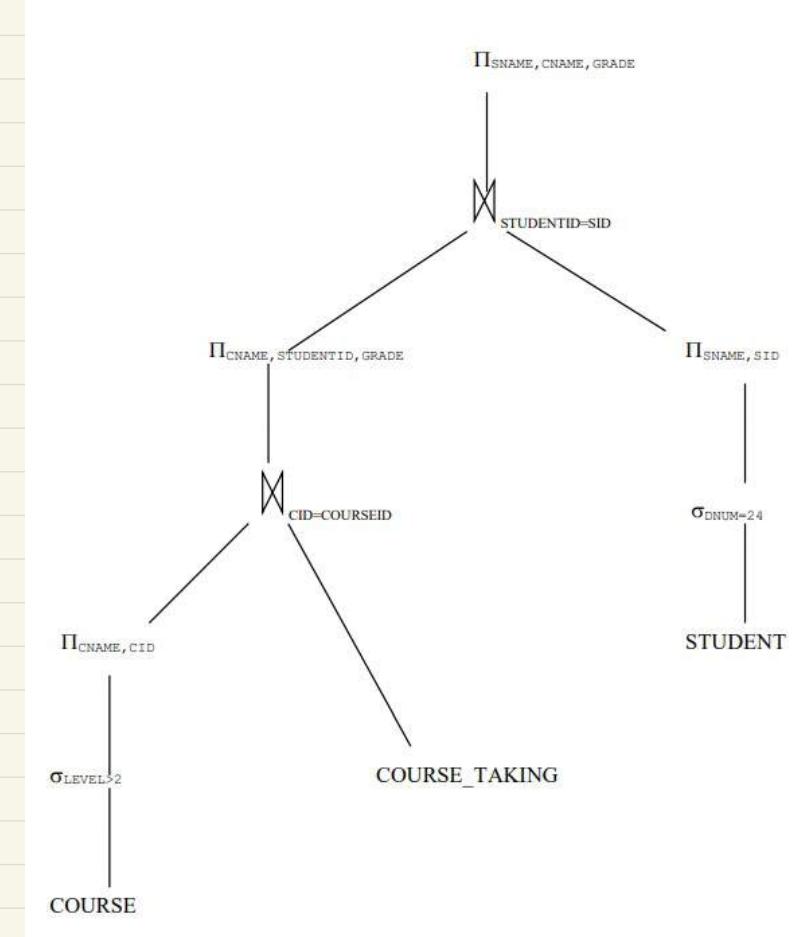
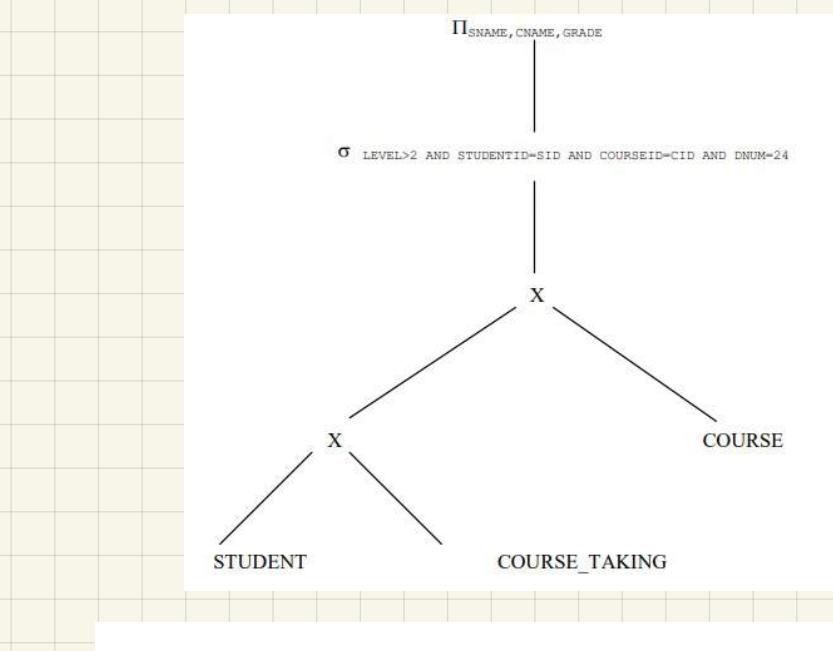
# Homework

## Example



**Figure 15.5** Steps in converting a query tree during heuristic optimization.

- (a) Initial (canonical) query tree for SQL query Q.
- (b) Moving SELECT operations down the query tree.
- (c) Applying the more restrictive SELECT operation first.
- (d) Replacing CARTESIAN PRODUCT and SELECT WITH JOIN operations.
- (e) Moving PROJECT operations down the query tree.



# Lecture 10-11 Transaction & Concurrency Control.

**Transaction boundary** logical unit of database processing that include 1 or more access operation  
begin - end

simplified model

- read-item(X)
  1. find the address of disk block that contain X
  2. copy disk block to buffer in memory
  3. copy item X from buffer to program variable named X.

- write-item(X)
  1. same
  2. same
  3. copy X from program variable name X into correct location in the buffer.
  4. store updated block from buffer back to disk

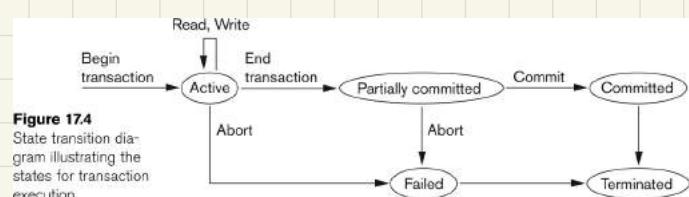
problem with transaction processing - concurrent execution is uncontrolled:

- lost update problem
- temporary update problem
- incorrect summary problem

why Transaction may fail

- computer failure
- transaction or system error
- local error or exception condition detected by transaction
- concurrency control enforcement
- disk failure
- physical problem and catastrophe

## Transaction States



Operation

- begin transaction
- read/ write
- end transaction
- commit transaction
- roll back

recovery technique

**undo** signal that operation end unsuccessfully operation must undone

**redo** certain transaction operation must redone to ensure all operations of transaction have been applied successfully to database

System log

ACID properties

- **ACID**
- atomicity (原子性) 要么不做要么做完
- consistency (一致性) 独立执行事务可以保持数据库的一致性
- isolation (隔离性) 每个事务必须不知道其他并发执行的事务
- durability (持久性)

transaction schedule interleaved fashion

recoverable schedule TB write  
TA read (TB written)  $\Rightarrow$  TB commit before TA

cascade rollback: failure lead to series roll back

T1	T2	T1	T2	T1	T2	T3
Read(A) Write(A)		Read(A) Write(A)		Read(A)		
	Read(A) Commit/Abort		Read(B) Commit/Abort		Read(A)	
Read(B) Commit/Abort			Commit/Abort			Commit/Abort

**Non-recoverable**  
If T2 commits and then  
T1 Aborts...?

**Recoverable**

**Recoverable but:**  
When T1 fails, T2 and T3 should rollback

cascadeless schedule committed before being read

strict schedule  $X$  read  $X$  write until last write operation  $T_n$  commit

Serial schedule consecutively.

Serializable schedule result equivalent produce same final state  
conflict equivalent order of conflicting operation same

Isolation level

## Concurrency - Two-phase locking techniques

lock(x) secure permission  
to read/to write data item for transaction

unlock(x) remove permission from data item

two lock mode  
- shared lock(x) read value x write }

- exclusive(x) write lock x read }

} no share can be applied in other T

wait for graph

two phase (mutually exclusive)

locking (growing)

unlocking (shrinking)

Dead lock prevention  $\rightarrow$  lock all data  
before begin execution